

```

1 import matplotlib.pyplot as plt
2 import numpy as np
3
4 # Define the functions and their derivatives
5 def f1(x): return x**2 # Convex function
6 def df1(x): return 2*x
7
8 def f2(x): return x**4 - 3*x**2 + 1 # Non-convex function
9 def df2(x): return 4*x**3 - 6*x
10
11 def gradient_descent(f, df, x0, eta, max_iter=10000, max_value=1e10):
12     iter = 0
13     x = x0
14     xs = [x0] # list to store x values
15     while iter < max_iter:
16         x_new = x - eta * df(x)
17         if abs(f(x_new)) > max_value:
18             print("Divergence!")
19             break
20         x = x_new
21         xs.append(x) # store x value
22         iter += 1
23     return xs, iter
24
25 x0 = 1
26
27 # Part (a)
28 eta = 0.01
29 xs, num_iter = gradient_descent(f1, df1, x0, eta)
30 print(f"Part (a) - For eta={eta}, minimum of f1: {xs[-1]:.6f} (found in {num_iter} iterations)")
31
32 # Plot f1
33 x = np.linspace(-1, 1, 400)
34 y = f1(x)
35 plt.plot(x, y, label="f1")

```

```

32 # Plot f1
33 x = np.linspace(-1, 1, 400)
34 y = f1(x)
35 plt.plot(x, y, label="f1")
36
37 # Plot path of gradient descent
38 plt.plot(xs, [f1(x) for x in xs], 'o-', label="Path of gradient descent")
39
40 # Add horizontal line at y=0
41 plt.axhline(0, color='r', linestyle='--', label="Minimum of f1")
42
43 plt.title("Convergence of gradient descent for small learning rate")
44 plt.xlabel("x")
45 plt.ylabel("f1(x)")
46 plt.legend()
47 plt.show()
48
49 # Part (b)
50 etas = [0.01, 0.1, 0.5, 1.0] # different values of eta
51 for eta in etas:
52     xs, num_iter = gradient_descent(f1, df1, x0, eta)
53     print(f"Part (b) - For eta={eta}, the method converged to the minimum of f1: {xs[-1]:.6f} in {num_iter} iterations")
54     plt.plot([f1(x) for x in xs], label=f"eta={eta}")
55
56 plt.title("Convergence of gradient descent for different learning rates")
57 plt.xlabel("Iteration")
58 plt.ylabel("f1(x)")
59 plt.legend()
60 plt.show()
61
62 # Part (c)
63 etas = [1.1, 1.5, 2.0] # different large values of eta
64 for eta in etas:

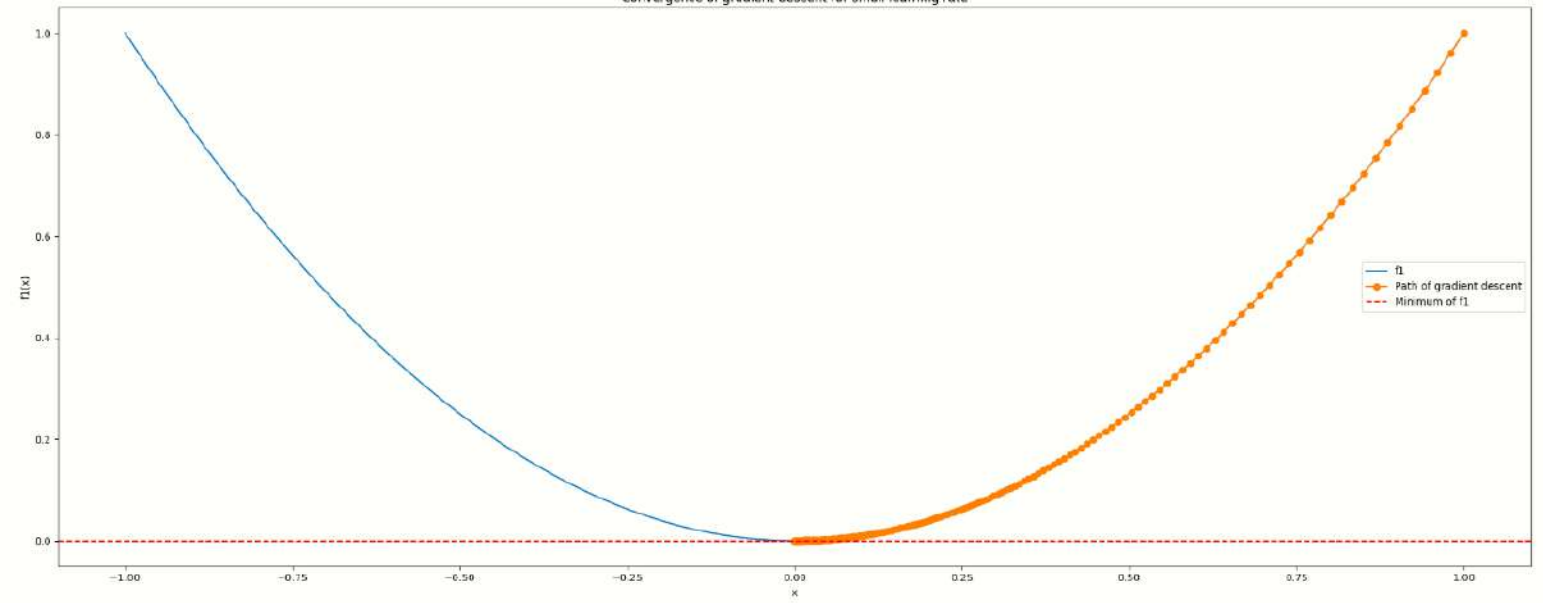
```

```

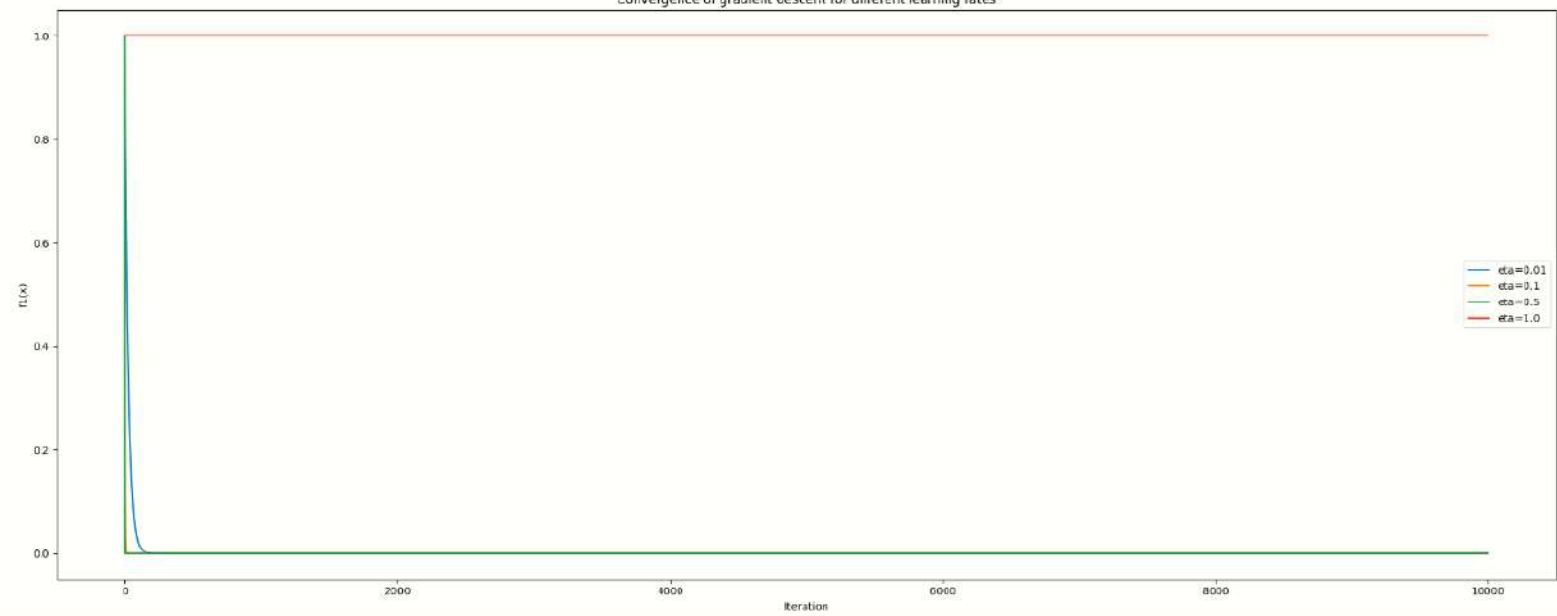
62 # Part (c)
63 etas = [1.1, 1.5, 2.0] # different large values of eta
64 for eta in etas:
65     xs, num_iter = gradient_descent(f1, df1, x0, eta)
66     if xs[-1] > 1e10:
67         print(f"Part (c) - For eta={eta}, the method diverged")
68     else:
69         print(f"Part (c) - For eta={eta}, the method converged to the minimum of f1: {xs[-1]:.6f} in {num_iter} iterations")
70     plt.plot([f1(x) for x in xs if f1(x) < 1e10], label=f"eta={eta}")
71
72 plt.title("Potential divergence of gradient descent for large learning rates")
73 plt.xlabel("Iteration")
74 plt.ylabel("f1(x)")
75 plt.xlim([0, 100]) # adjust x-axis range
76 plt.ylim([0, 100]) # adjust y-axis range
77 plt.legend()
78 plt.show()
79
80 # Part (d)
81 eta = 0.01
82 xs, num_iter = gradient_descent(f2, df2, x0, eta)
83 print(f"Part (d) - For eta={eta}, minimum of f2: {xs[-1]:.6f} (found in {num_iter} iterations)")
84
85 # Plot f2
86 x = np.linspace(-2, 2, 400)
87 y = f2(x)
88 plt.plot(x, y, label="f2")
89
90 # Plot path of gradient descent
91 plt.plot(xs, [f2(x) for x in xs], 'o-', label="Path of gradient descent")
92
93 plt.title("Gradient descent can get stuck in a local minimum for non-convex functions")
94 plt.xlabel("x")
95 plt.ylabel("f2(x) = x^4 - 3x^2 + 1")
96 plt.legend()

```

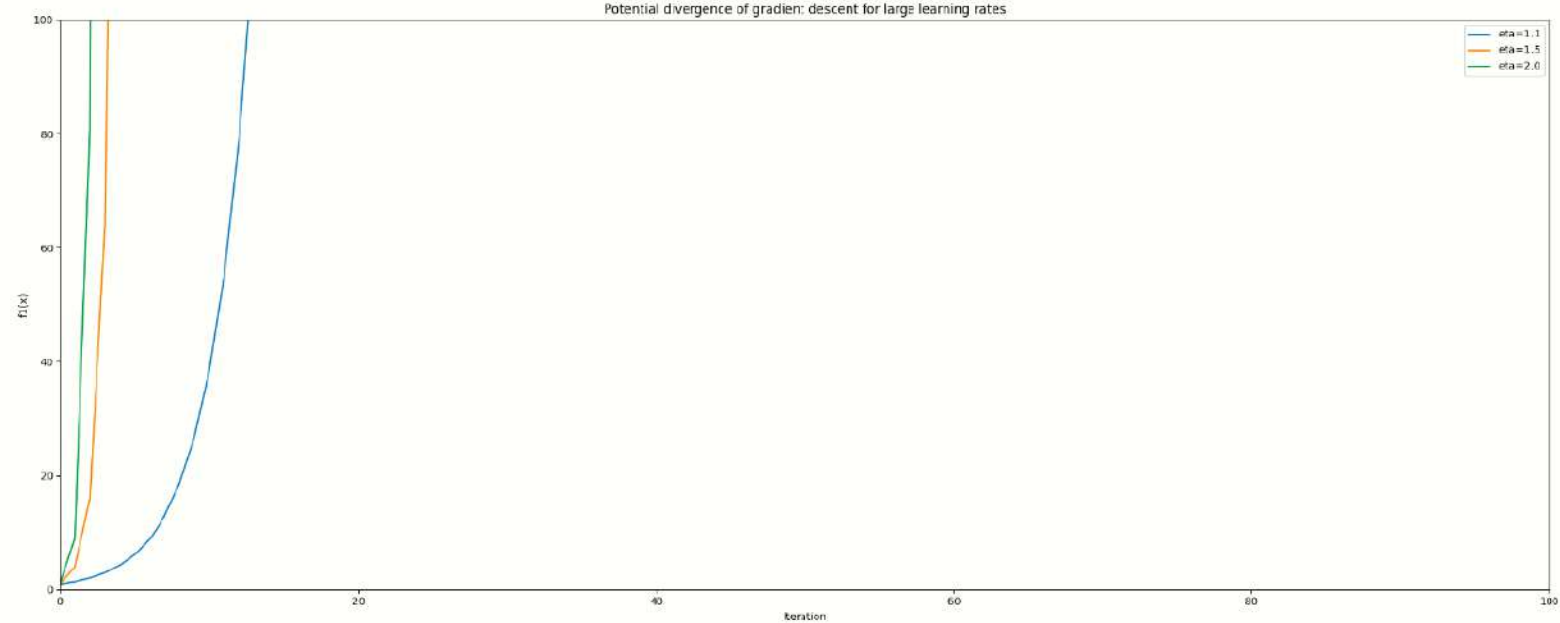
Convergence of gradient descent for small learning rate

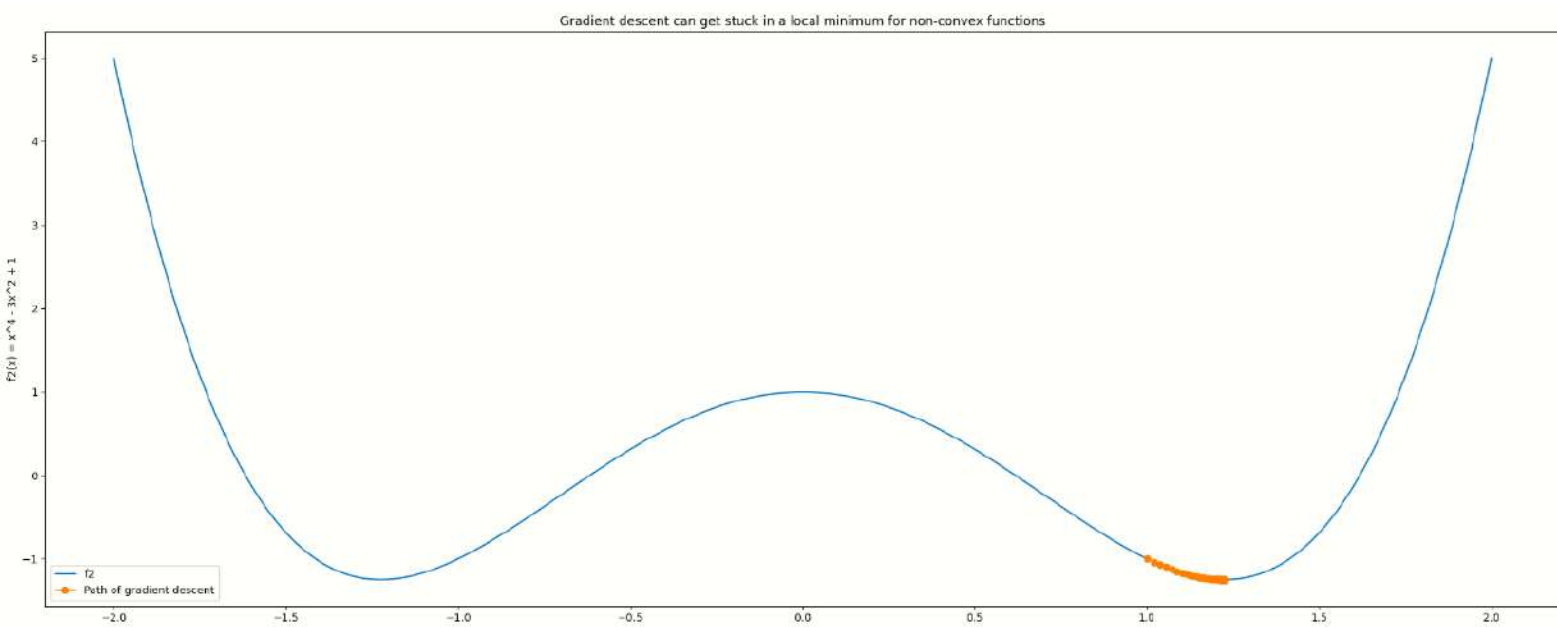


Convergence of gradient descent for different learning rates



Potential divergence of gradient descent for large learning rates





Part (a) - For  $\eta=0.01$ , minimum of  $f_1$ : 0.000000 (found in 10000 iterations)  
Part (b) - For  $\eta=0.01$ , the method converged to the minimum of  $f_1$ : 0.000000 in 10000 iterations  
Part (b) - For  $\eta=0.1$ , the method converged to the minimum of  $f_1$ : 0.000000 in 10000 iterations  
Part (b) - For  $\eta=0.5$ , the method converged to the minimum of  $f_1$ : 0.000000 in 10000 iterations  
Part (b) - For  $\eta=1.0$ , the method converged to the minimum of  $f_1$ : 1.000000 in 10000 iterations  
Divergence!  
Part (c) - For  $\eta=1.1$ , the method converged to the minimum of  $f_1$ : -97368.504802 in 63 iterations  
Divergence!  
Part (c) - For  $\eta=1.5$ , the method converged to the minimum of  $f_1$ : 65536.000000 in 16 iterations  
Divergence!  
Part (c) - For  $\eta=2.0$ , the method converged to the minimum of  $f_1$ : 59049.000000 in 10 iterations  
Part (d) - For  $\eta=0.01$ , minimum of  $f_2$ : 1.224745 (found in 10000 iterations)